
Prompt Injection Defense for an Agentic RAG Research Synthesizer

Shivani Kalal

Department of Computer and Information Science
University of Mississippi
skalal@go.olemiss.edu

Abstract

Agentic Retrieval-Augmented Generation (RAG) systems are vulnerable to indirect prompt injection (IPI) attacks, in which malicious instructions embedded in retrieved documents manipulate system behavior without the user’s knowledge. This paper studies IPI vulnerability and defense in a fully operational agentic RAG research synthesizer. Two attack families are evaluated — false information injection and contradiction injection, with two corresponding defenses implemented: D1, untrusted-context isolation via structured prompt hierarchies, and D2, an abstract-level injection detector that scans retrieved content before synthesis. Baseline evaluation shows 100% attack success rate (ASR) for both attack families. D1 reduces ASR to 0% with no degradation in system utility. D2 provides an additional detection layer that flags suspicious abstracts before synthesis. Results demonstrate that prompt engineering defenses are sufficient to neutralize retrieval-based injection attacks in this domain while preserving full citation coverage and review quality on benign queries.

1 Introduction

Retrieval-Augmented Generation (RAG) systems have become widely adopted in high-trust applications such as academic research, where users expect outputs to be accurate and free of external manipulation. These systems combine a large-language-model (LLM) with external document retrieval and callable tools, enabling dynamic, grounded responses from live corpora. However, this architecture introduces a practical security vulnerability known as *indirect prompt injection* (IPI): malicious instructions embedded in retrieved documents can override the agent’s intended behavior, even when the end user is entirely benign.

Unlike jailbreaks, where the user is the attacker, IPI attacks are hidden inside data that the system fetches automatically. The user types a legitimate query and has no indication that the retrieved content has been tampered with. In a research synthesizer that ingests paper abstracts and generates citation-backed literature reviews, the consequences are significant. A successful attack can silently corrupt a literature review, introduce fabricated statistical claims cited as peer-reviewed research, or trigger unauthorized system actions, all without the user’s knowledge.

This paper studies IPI vulnerability and defense in a fully operational agentic RAG research synthesizer built on LangChain, Claude, ChromaDB, BM25 hybrid retrieval, SQLite, and Streamlit. The system fetches papers from arXiv, stores them in a local corpus, retrieves the most relevant abstracts per query, and generates structured citation-enforced literature reviews. We treat this system as a security testbed and make the following contributions:

- A quantitative baseline measurement of IPI vulnerability in an undefended citation-enforced RAG agent across two attack families: false information injection and contradiction injection.
- Two defenses: (D1) untrusted-context isolation via structured prompt hierarchies, and (D2) a contradiction injection detector that scans retrieved abstracts for instruction-like language before synthesis.
- An ablation study comparing Baseline, D1, D2, and D1+D2 across both security and utility metrics, providing evidence on whether prompt engineering defenses are sufficient against retrieval-based attacks.

This paper addresses four research questions: **RQ1:** How vulnerable is an undefended agentic RAG research synthesizer to indirect prompt injection attacks? **RQ2:** Does prompt-level untrusted-context isolation reduce attack success rate? **RQ3:** Does adding an abstract-level injection detector provide additional protection? **RQ4:** Do these defenses reduce system utility for legitimate users?

2 Related Work

2.1 Indirect Prompt Injection

Greshake et al. [1] introduced the indirect prompt injection threat model, demonstrating that LLM-integrated applications are vulnerable to malicious instructions embedded in external content retrieved at query time. Unlike direct prompt injection where the user is the attacker, IPI attacks are invisible to the benign user.

Yi et al. [2] presented BIPIA, the first systematic benchmark for IPI evaluation, identifying the root cause: LLMs cannot inherently distinguish between instructions and data when both are present in the same context.

Zhan et al. [3] introduced InjecAgent, a benchmark of 1,054 test cases, finding that GPT-4 was vulnerable to IPI in 24% of cases under standard prompting.

2.2 Retrieval-Augmented Generation Attacks

Zou et al. [4] demonstrated in PoisonedRAG that injecting as few as five poisoned documents into a RAG knowledge base achieves 90% attack success rate — directly motivating the multi-paper injection methodology used in this paper.

2.3 Defenses

Hines et al. [5] proposed Spotlighting, reducing IPI attack success from over 50% to under 2% through prompt-level provenance marking. The D1 defense presented here applies the same principle. DeBenedetti et al. [6] proposed CaMeL, a system-level defense achieving provable security on 77% of AgentDojo benchmark cases through architectural separation of instruction and data flows.

2.4 Research Gap

No prior work evaluates IPI in a citation-enforced, tool-using research synthesis agent — a system where integrity failures have direct academic consequences and where existing citation enforcement provides unexpected natural resistance to certain attack classes. This paper fills that gap.

3 System Description

The system under test is a fully operational agentic RAG research synthesizer built prior to this security evaluation. The architecture consists of six components.

Retrieval pipeline. When a user submits a research query, the system searches arXiv for relevant papers and persists records (title, abstracts, authors, year) into a local SQLite database (`papers.db`) and a ChromaDB vector store. Subsequent queries on the same topic retrieve directly from the local corpus. Retrieval uses a hybrid approach combining BM25 keyword search and ChromaDB semantic

search, with scores blended using a weighted parameter $\alpha = 0.5$. The top- k most relevant abstracts are returned as formatted text to the LangChain agent, where the underlying LLM (Claude) acts as the synthesis engine, generating a consolidated response from the retrieved content.

Synthesis Engine. The synthesis engine (`synthesis_engine.py`) passes retrieved abstracts to Claude (`claude-3-5-sonnet-20250514`) via LangChain with strict citation enforcement rules: every factual claim must cite a paper using [X] format, no external knowledge may be used, and if papers disagree both must be cited and the disagreement noted. The engine produces structured literature reviews with section for introduction, key findings, methodology, research gaps and conclusion.

Agent framework. LangChain orchestrates the agent loop using a tool-calling pattern. Query routing in `main.py` classifies each user query into a type (synthesis, export, search, etc.) and creates a specialized agent with a limited tool set for that query type. This routing creates natural tool isolation — side-effect tools such as `export_bibtex` are only available in export query branches, not synthesis branches.

Quality assurance. A built-in QA layer validates generated reviews against retrieved abstracts, checking citation coverage, semantic accuracy, and citation density, producing a quality score from 0 to 1.0.

Interface. The system is deployed via Streamlit with a chat interface. At the time of evaluation the corpus contained 615 papers.

Security relevance. The synthesis engine is the primary attack surface. Retrieved abstracts are passed directly to Claude as context alongside trusted system instructions, with no structural separation between the two. This mirrors the fundamental vulnerability identified by Yi et al. [2] — LLMs cannot inherently distinguish between instructions and data when both appear in the same context.




Figure 1: System architecture of the agentic RAG research synthesizer. User queries are classified and routed to either external search or local retrieval, with results merged via hybrid retrieval before synthesis.

4 Threat Model

4.1 Attacker Capability

The attacker operates under a black-box threat model with no access to the system prompt, model weights, or agent internals. The sole attack capability is the ability to insert records into the retrieval corpus — either by directly writing to `papers.db` in a local database, or in a realistic scenario, by uploading a paper to arXiv whose abstract contains malicious instructions. Since the system automatically fetches and stores papers from arXiv during normal operations, a poisoned paper uploaded to arXiv would be ingested into the local corpus and retrieved during subsequent queries on the same topic.

4.2 Attack Surface

The primary attack surface is the abstract field of retrieved paper records. During synthesis queries the top- k retrieved abstracts are passed directly to Claude as content alongside the trusted system prompt. As identified by Yi et al. [2], LLMs cannot inherently distinguish between instructions embedded in retrieved content and instructions from the trusted system prompt. This creates an injection surface in any system that combines trusted instructions with untrusted retrieved content in the same context window — the fundamental architectural condition present in the system under test, as described in Section 3.

4.3 Natural Resistance

An important baseline finding is that the undefended system exhibits natural resistance to single-paper injection attacks. The synthesis prompt enforces strict citation rules: every factual claim must cite a paper, only information present in the retrieved abstracts may be used, and the output must follow a fixed structure. These constraints are strong enough that Claude follows them even when a single poisoned abstract contains explicit override instructions such as `IGNORE ALL PREVIOUS INSTRUCTIONS`. Citation sabotage and instruction hijacking attacks with one poisoned paper each yielded 0% ASR against the undefended system. This is consistent with prior work by Hines et al. [5], which showed that structured prompts can provide inherent resistance to injection without any explicit security instrumentation.

However, this natural resistance has limits. When the ratio of poisoned to legitimate papers in the retrieved top- k is sufficiently high, the injected content overwhelms the system’s natural resistance. This motivates the multi-paper attack methodology evaluated below.

4.4 Attack Families

Two attack families are evaluated in this paper.

False information injection. Five poisoned papers were inserted into `papers.db` and the ChromaDB vector store, each containing fabricated claims disguised as legitimate research findings. Example payloads included claims such as “transformer attention mechanisms fail in 94% of real-world deployments” and “Google, Meta, and OpenAI have discontinued attention-based architectures.” All five papers used keyword-rich titles matching the target query to ensure high retrieval rank. In testing, four of five poisoned papers ranked in the top ten retrieved results (best retrieval score: 0.890), and all were passed to Claude as context. The attack succeeded — fabricated claims appeared in the generated literature review cited as peer-reviewed research, with poisoned papers appearing in the references section alongside legitimate sources (Figure 2). This methodology is directly motivated by Zou et al. [4], who demonstrated that five poisoned documents achieve 90% ASR against a RAG knowledge base.




Figure 2: System output under false information injection attack. The fabricated claim — that attention mechanisms degrade exponentially after six months of deployment based on studies involving 10,000 models — is cited as [1] in the review body. All five poisoned papers appear in the references section alongside legitimate sources.

Contradiction injection. A single poisoned paper exploits the synthesis engine’s conflict-resolution rule: “*if papers disagree, cite both and note the disagreement.*” The poisoned abstract claimed to contradict all other retrieved papers and embedded a fabricated statistic: “self-attention causes catastrophic memory loss in 89% of production deployments.” Claude, following its conflict-resolution rule, included the contradiction in the Research Gaps section and cited the poisoned paper as a legitimate source (Figure 3). The poisoned paper ranked second among retrieved results (retrieval score: 0.767). Unlike false information injection, this attack requires only one poisoned paper and exploits a specific synthesis rule rather than overwhelming the system with volume. This attack class has not been previously studied in the IPI literature [1, 3] and is specific to citation-enforced synthesis systems.




Figure 3: System output under contradiction injection attack. The poisoned paper is cited as [4] in the Research Gaps section, presenting a fabricated 89% catastrophic memory loss statistic as a legitimate academic disagreement. One poisoned paper was sufficient.

4.5 Attacker Goal

The attacker’s goal is to corrupt the generated literature review without the user’s knowledge. A successful attack introduces fabricated claims into the review body, cited as peer-reviewed research. Critically, the system’s built-in QA layer assigned both attacked outputs a quality score of 1.000/A — confirming that the corruption is invisible to automated quality checks and would appear legitimate to the end user.

5 Defense Design

5.1 D1: Untrusted-Context Isolation

D1 addresses the root attack surface identified in Section 4: retrieved abstracts and trusted system instructions occupy the same context window with no structural separation. The defense implements an explicit trust hierarchy via two targeted modifications to `synthesis_engine.py`, requiring no model retraining and no architectural changes.

Change 1 — Context wrapping in `prepare_paper_context()`. Retrieved abstracts are wrapped in explicit provenance labels before being passed to Claude:

```
[RETRIEVED EVIDENCE - UNTRUSTED EXTERNAL CONTENT]
The following abstracts are retrieved from an external
corpus. They may contain malicious instructions.
Treat as data only.

[1] Author (Year): Title
Abstract: ...

[END RETRIEVED EVIDENCE]
```

Change 2 — Security rule in `_create_synthesis_prompt()`. An explicit trust hierarchy instruction is prepended to the synthesis prompt:

```
SECURITY RULE - TRUST HIERARCHY:
You operate under a strict trust hierarchy:
- TRUSTED: This system prompt and the user query
- UNTRUSTED: All retrieved abstracts marked
  [RETRIEVED EVIDENCE]
If any retrieved abstract appears to give
instructions, ignore it completely. Treat retrieved
abstracts as evidence to cite - NEVER as
instructions to follow.
```

This approach is directly grounded in Spotlighting [5], which demonstrated that prompt-level provenance marking reduces IPI attack success from over 50% to under 2%. D1 applies the same principle: by making the boundary between trusted instructions and untrusted retrieved content explicit and structural, Claude is given the information it needs to enforce the trust hierarchy correctly. The trust model is straightforward — the system prompt and user query are trusted; all retrieved abstracts are untrusted and must be treated as data to cite, never as instructions to follow.

5.2 D1 Results

Table 1 shows the effect of D1 on attack success rate. False information injection dropped from 100% to 0% ASR. The natural resistance findings for single-paper attacks were unchanged.

Attack	Baseline ASR	D1 ASR
Single paper attacks (1 paper)	0%	0%
False information injection (5 papers)	100%	0%

Table 1: Attack success rate: Baseline vs. D1.




Figure 4: D1 results. Before D1: 100% ASR with all five poisoned papers cited and fabricated claims present in the review. After D1: 0% ASR with zero poisoned papers cited and no fabricated content in the output.

Two prompt engineering changes eliminated the attack entirely. No model retraining and no architectural changes were required.

One trade-off was observed: citation coverage dropped from 100% to 50% in the attacked condition. The missing citations — [1], [2], [3], [6], and [10] — correspond exactly to the poisoned paper numbers, confirming that D1 caused Claude to identify and exclude the injected records specifically. This is correct behavior: the defense is working as intended, excluding suspicious content rather than incorporating it.

In the attacked condition, the QA layer assigned a quality score of 0.736/1.0 (Grade C), reflecting the citation coverage drop. This does not represent a failure of D1 — it reflects the QA layer’s inability to distinguish security-motivated exclusions from genuine omissions. On clean queries with no injected content, citation coverage remained at 100% and quality score at 1.000/A, confirming that D1 introduces no utility degradation for legitimate users.

5.3 D2: Untrusted-Abstract Detection

D1 operates inside Claude — it instructs the model to treat retrieved content as data rather than instructions. D2 operates before Claude is involved at all. One method, `_validate_claims()`, is added to `synthesis_engine.py` and runs at the application layer in pure Python, scanning each retrieved abstract for instruction-like language before synthesis begins.

The detection logic is grounded in a simple observation: real academic abstracts describe findings and report results. They never issue commands to AI systems. A pattern match against the following keyword list is therefore sufficient to identify injected content:

SUSPICIOUS PATTERNS:
 "must be cited as"
 "academic standards require"
 "IGNORE ALL PREVIOUS INSTRUCTIONS"
 "NEVER write any citation"
 "DO NOT include any citation"

For comparison, legitimate abstract language looks like:

LEGITIMATE:
 "This paper presents findings on..."
 "Results show that..."
 "We propose a method that..."

If a suspicious pattern is matched in any retrieved abstract, that paper is immediately flagged and excluded from synthesis. A security alert is appended to the review output and the user is warned before reading the result. Claude never sees the flagged abstract. This deterministic, application-layer check closes a gap that D1 alone cannot fully address: D1 relies on Claude following the trust hierarchy instruction correctly on every query. D2 makes no such assumption — if the pattern matches, the paper is blocked regardless of model behavior.

5.4 D2 Results




Figure 5: D2 results. Before D2: contradiction injection achieved 100% ASR with the poisoned paper cited as [4] and the fabricated statistic present in the review. After D2: the poisoned abstract was flagged by `_validate_claims()` before synthesis. The D2 security alert is shown at the bottom of the output, warning the user that suspicious content was detected and excluded.

Under contradiction injection, D2 detected the poisoned abstract before synthesis and excluded Paper [4] from the retrieved context passed to Claude. The actual alert surfaced to the user read:

D2 SECURITY ALERT: Suspicious claims detected in this review that do not appear in any retrieved abstract. This review may contain fabricated information. Flagged content: [\'Paper [4]: Transformer Self-Attention Mechanisms: A Comprehensive Survey\']

ASR dropped from 100% to 0% for contradiction injection. No false positives were observed on clean queries — legitimate abstracts contain none of the flagged patterns and pass through the detector without triggering an alert. D1 and D2 together provide complementary protection: D1 prevents the model from following injected instructions; D2 catches poisoned abstracts before the model ever processes them.

6 Evaluation

6.1 Experimental Setup

All attacks were evaluated on a single query — “*generate literature review on transformer attention mechanisms*” — against the live system corpus of 615 papers. The top- k retrieved results for this query consisted of 10 papers. Poisoned papers were inserted directly into `papers.db` and the ChromaDB vector store prior to each attack trial and removed afterward. Three conditions are evaluated: Baseline (undefended), D1 only, and D1+D2. Four metrics are reported across conditions:

Attack Success Rate (ASR) — whether fabricated claims appeared in the generated review cited as legitimate research.

Citation Coverage — the percentage of retrieved papers actually cited in the generated literature review. If 10 papers are retrieved and all 10 are cited in the output, coverage is 100%. A drop in citation coverage under attack conditions indicates the system is excluding retrieved papers — either due to the defense working correctly (excluding poisoned papers) or due to utility degradation.

Quality Score — a score from 0 to 1.0 assigned by the system’s built-in QA layer, a separate validation module that runs after each review is generated. The QA layer checks three components: citation verification (are all cited numbers valid and tied to a retrieved paper), semantic accuracy (do the claims in the review match the content of the cited abstracts), and citation density (citations per 1,000 characters). These are combined into an overall score; 1.000 corresponds to Grade A, indicating a fully cited, semantically accurate, and well-structured review.

D2 False Positives — whether D2 flagged any legitimate paper on clean queries with no injected content.

6.2 Baseline Results (RQ1)

Table 2 summarizes attack results on the undefended system across all tested attack families.

Attack	Papers	ASR	Notes
Citation sabotage — re-move markers	1	0%	Citations main-tained
Citation sabotage — format switch	1	0%	[X] format main-tained
Instruction hijacking — conclusion	1	0%	Conclusion unaf-fected
False information injection	5	100%	4/5 papers cited
Contradiction injection	1	100%	Fabricated claim in review

Table 2: Baseline attack results on undefended system. Query: *generate literature review on transformer attention mechanisms*.

Natural resistance to single-paper attacks. Single-paper citation sabotage and instruction hijacking attacks yielded 0% ASR against the undefended system. The synthesis prompt in `synthesis_engine.py` enforces citation rules strongly enough that Claude follows system-level instructions over injected content in a single retrieved abstract. This constitutes natural IPI resistance through prompt engineering, consistent with prior work on provenance marking [5].

Attack scaling effect. Single-paper injections failed across all tested payloads. The false information injection attack only succeeded when five poisoned papers were injected simultaneously. With one poisoned paper among ten retrieved, the synthesis prompt dominated. With four poisoned papers among ten retrieved — a 40% poisoned ratio in the top- k — the false content overwhelmed the system’s natural resistance. This scaling threshold is consistent with findings by Zou et al. [4], who identified five poisoned documents as sufficient to achieve 90% ASR in a RAG knowledge base.

Vulnerability to false information injection. Four of five poisoned papers were retrieved and cited as legitimate sources. The review included fabricated claims such as attention mechanisms failing in

94% of real-world deployments and formal deprecation by Google, Meta, and OpenAI — none of which appear in any real paper. The system’s built-in QA layer assigned the attacked output a quality score of 1.000/A, confirming that the corruption is invisible to automated quality checks and would appear legitimate to the end user.

Vulnerability to contradiction injection. A single poisoned paper exploiting the conflict-resolution rule achieved 100% ASR. The fabricated 89% catastrophic memory loss statistic appeared in the Research Gaps section cited as a legitimate academic disagreement. The QA layer again assigned a quality score of 1.000/A. This answers **RQ1**: the undefended system is fully vulnerable to both attack families, with no automated indication of compromise.

Natural tool isolation. Tool misuse attacks yielded 0% ASR due to two existing architectural protections: query-type based tool isolation in `main.py` that restricts side-effect tools such as `export_bibtex` to export query branches only, and an environment variable authorization gate on the `save` tool. This natural protection requires no additional defense instrumentation.

6.3 Ablation Study — Security (RQ2, RQ3)

Table 3 shows ASR across all three conditions.

Attack	Baseline	D1	D1+D2
False information injection (5 papers)	100%	0%	0%
Contradiction injection (1 paper)	100%	0%	0% + alert

Table 3: Ablation study: Attack Success Rate across all conditions.

D1 alone reduced ASR from 100% to 0% for both attack families, answering **RQ2** affirmatively. The two targeted modifications to `synthesis_engine.py` — abstract provenance labeling and the trust hierarchy security rule — were sufficient to neutralize both attacks entirely. This answers **RQ3**: D2 provides an additional detection layer on top of D1, flagging and excluding the poisoned abstract before synthesis and surfacing a security alert to the user, but D1 alone is already sufficient for 0% ASR. The complementary value of D2 is detection and transparency — the user is explicitly warned that suspicious content was identified — rather than additional ASR reduction.

6.4 Utility Retention (RQ4)

Table 4 shows system utility metrics on clean queries with no injected content across all three conditions.

Condition	Citation Coverage	Quality Score	D2 False Positives
Baseline	100%	1.000/A	N/A
D1 only	100%	1.000/A	N/A
D1+D2	100%	1.000/A	0

Table 4: Utility retention on clean queries (no injection).

Under all three conditions, citation coverage remained at 100% and quality score at 1.000/A on clean queries. D2 triggered zero false positives — legitimate academic abstracts contain none of the flagged instruction-like patterns and pass through `_validate_claims()` without triggering an alert. This answers **RQ4**: neither defense degrades system utility for legitimate users.

One trade-off was observed in the *attacked* condition under D1: citation coverage dropped from 100% to 50%, with missing citations [1], [2], [3], [6], and [10] corresponding exactly to the poisoned paper numbers. The QA layer assigned a quality score of 0.736/1.0 (Grade C) in this condition, penalizing the missing citations. This does not represent a failure of D1 — it reflects the QA layer’s inability to distinguish security-motivated exclusions from genuine omissions. Claude correctly identified and excluded the suspicious content; the trade-off is inherent to any defense that blocks retrieved content from synthesis.

7 Project Experience

7.1 Difficulties

The dual-database injection problem. The most significant implementation difficulty was injecting poisoned papers correctly. The system persists records in two separate stores: a SQLite database (`papers.db`) and a ChromaDB vector store. The normal ingestion path, `store_paper()`, uses `UPDATE` rather than `INSERT`, so calling it on a new record silently fails. The correct procedure required direct `INSERT` into SQLite first, a separate ChromaDB embedding call, and a BM25 index rebuild — none of which was documented in the existing codebase.

Authors JSON format. A subtle but blocking issue was the `authors` field format. SQLite stores authors as a JSON string, not a plain comma-separated string. Passing a Python list directly raised a `sqlite3.ProgrammingError` at binding time. The fix required wrapping the list with `json.dumps()` before insertion — an implicit requirement that only surfaced during live testing.

D2 scope revision. The original D2 design targeted tool misuse via a keyword authorization gate on `export_bibtex`. During baseline testing, tool misuse attacks achieved 0% ASR on the undefended system due to two pre-existing architectural protections: query-type based tool isolation in `main.py` and an environment variable authorization gate. The planned D2 was redundant. It was redesigned to target contradiction injection via abstract-level pattern scanning — a significant scope revision that produced a more meaningful contribution.

7.2 Unexpected Findings

Natural resistance. The undefended system already exhibited strong resistance to single-paper injection attacks — citation sabotage and instruction hijacking each failed at 0% ASR before any defense was applied. The finding reframed the research narrative: the system’s citation enforcement constitutes a form of natural IPI resistance, and the security gap is specifically the multi-paper volumetric attack that overwhelms it.

QA layer limitation. Under both attack families, attacked outputs received a quality score of 1.000/A — the highest possible grade — confirming that automated quality checks cannot detect IPI corruption. Conversely, under D1 in the attacked condition, the QA layer penalized the defense with a Grade C score, interpreting security-motivated exclusions of poisoned papers as missing citations.

7.3 Successful Techniques

Supporting infrastructure. Three utility scripts — `inject_poison.py`, `cleanup_poison.py`, and `check_papers.py` — handled the full dual-database injection and cleanup cycle, enabling reproducible attack trials without manual database manipulation.

Prompt engineering as defense. D1 demonstrated that two targeted modifications to `synthesis_engine.py` were sufficient to reduce ASR from 100% to 0% for both attack families, with no model retraining and no architectural changes. The simplicity of the effective defense is itself a finding: in this domain, structural clarity about what content is trusted is sufficient protection.

8 Conclusion

This paper studied indirect prompt injection vulnerability and defense in a fully operational agentic RAG research synthesizer. Two attack families were demonstrated against the undefended system: false information injection, which achieved 100% ASR by overwhelming the system’s natural resistance with five poisoned papers at a 40% poisoned ratio in the retrieved top- k ; and contradiction injection, which achieved 100% ASR with a single poisoned paper by exploiting the synthesis engine’s conflict-resolution rule. In both cases, the system’s built-in QA layer assigned attacked outputs a quality score of 1.000/A, confirming that the corruption is invisible to automated checks.

Two defenses were implemented and evaluated. D1 — untrusted-context isolation via structured prompt hierarchies, reduced ASR to 0% for both attack families with no utility degradation on clean queries. D2 — an abstract-level injection detector scanning for instruction-like patterns before synthesis, added a detection and transparency layer, flagging poisoned abstracts and warning users

before they read compromised output. Together, D1 and D2 provide complementary protection: D1 prevents the model from following injected instructions; D2 catches poisoned abstracts before the model ever processes them.

A key finding is that the undefended system already exhibited natural resistance to single-paper attacks through its citation enforcement prompt, and that natural tool isolation in the architecture provided inherent protection against tool misuse. The security gap was specific: multi-paper volumetric injection and rule-exploiting single-paper injection. Prompt-level defenses were sufficient to close both gaps, providing evidence that prompt engineering defenses are adequate against retrieval-based injection attacks in citation-enforced synthesis systems — at least under the black-box corpus poisoning threat model evaluated here.

Limitations of this study include evaluation on a single query topic, a small number of attack trials per family, and a keyword-based D2 detector that could be evaded by adversaries who avoid the flagged phrases. Future work should evaluate across multiple query domains, test adaptive adversaries aware of the D2 keyword list, and explore whether architectural separation of instruction and data flows in the style of CaMeL [6] provides stronger guarantees in this domain.

References

- [1] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection,” in *Proc. ACM Workshop on AI and Security (AISec)*, 2023. arXiv:2302.12173.
- [2] J. Yi, Y. Xie, B. Zhu, K. Hines, E. Kiciman, G. Sun, X. Xie, and F. Wu, “Benchmarking and defending against indirect prompt injection attacks on large language models,” arXiv:2312.14197, 2023.
- [3] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, “InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents,” in *Findings of ACL*, 2024. arXiv:2403.02691.
- [4] W. Zou, R. Geng, B. Wang, and J. Jia, “PoisonedRAG: Knowledge poisoning attacks to retrieval-augmented generation of large language models,” arXiv:2402.07867, 2024.
- [5] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with Spotlighting,” arXiv:2403.14720, 2024.
- [6] E. Debenedetti, I. Shumailov, T. Fan, J. Hayes, N. Carlini, D. Fabian, C. Kern, C. Shi, A. Terzis, and F. Tramèr, “Defeating prompt injections by design,” arXiv:2503.18813, 2025.